Low-Light Image Enhancement Using MIRNet

Hitesh Shanmukha B22AI013

Abstract

Enhancing low-light images is a challenging task in the field of image processing due to noise, color distortion, and loss of details. This report explores the MIRNet model, a multi-scale residual network designed to enhance low-light images by effectively capturing global context and local details. The model architecture combines multi-scale feature extraction, selective feature fusion, and attention mechanisms to produce high-quality enhanced images. Quantitative and qualitative results demonstrate the effectiveness of MIRNet in enhancing low-light images.

1 Introduction

Low-light conditions can significantly degrade the visual quality of images, affecting applications in photography, security surveillance, medical imaging, and more. Traditional image enhancement techniques often struggle to balance noise reduction and detail preservation. Convolutional Neural Networks (CNNs) have shown promise in this domain, but designing a network that can handle the complexities of low-light enhancement requires careful consideration.

This report focuses on the **MIRNet** (Multi-Scale Residual Network) model for low-light image enhancement. MIRNet learns enriched feature representations by combining contextual information from multiple scales while preserving high-resolution spatial details.

1.1 Background and Motivation

The challenges in low-light image enhancement include:

- **Noise Amplification**: Low-light images have a low signal-to-noise ratio.
- **Color Distortion**: Inadequate lighting affects color fidelity.
- **Detail Loss**: Fine details are often lost due to underexposure.

MIRNet addresses these challenges by integrating multi-scale feature extraction and attention mechanisms within a recursive residual framework.

2 Intuition Behind MIRNet

The core idea of MIRNet is to integrate information from multiple resolutions while maintaining the original high-resolution features. This approach allows the network to understand the image at different scales, capturing both global structures and fine details.

2.1 Key Components

1. **Multi-Scale Feature Extraction**: Processes the image at multiple scales to capture contextual information.
2. **Selective Kernel Feature Fusion (SKFF)**: Dynamically fuses features from different scales.
3. **Dual Attention Mechanisms**: Applies channel and spatial attention to emphasize relevant features.
4. **Recursive Residual Design**: Simplifies learning by progressively breaking down the input signal.

3 Model Architecture

3.1 Overview

The MIRNet architecture consists of several key components that work synergistically:

- **Multi-Scale Residual Blocks (MSRBs)**: Capture features at multiple scales.
- **Dual Attention Units (DAUs)**: Refine features using attention mechanisms.
- **Selective Kernel Feature Fusion (SKFF)**: Fuse multi-scale features dynamically.
- **Recursive Residual Groups (RRGs)**: Stack MSRBs recursively.

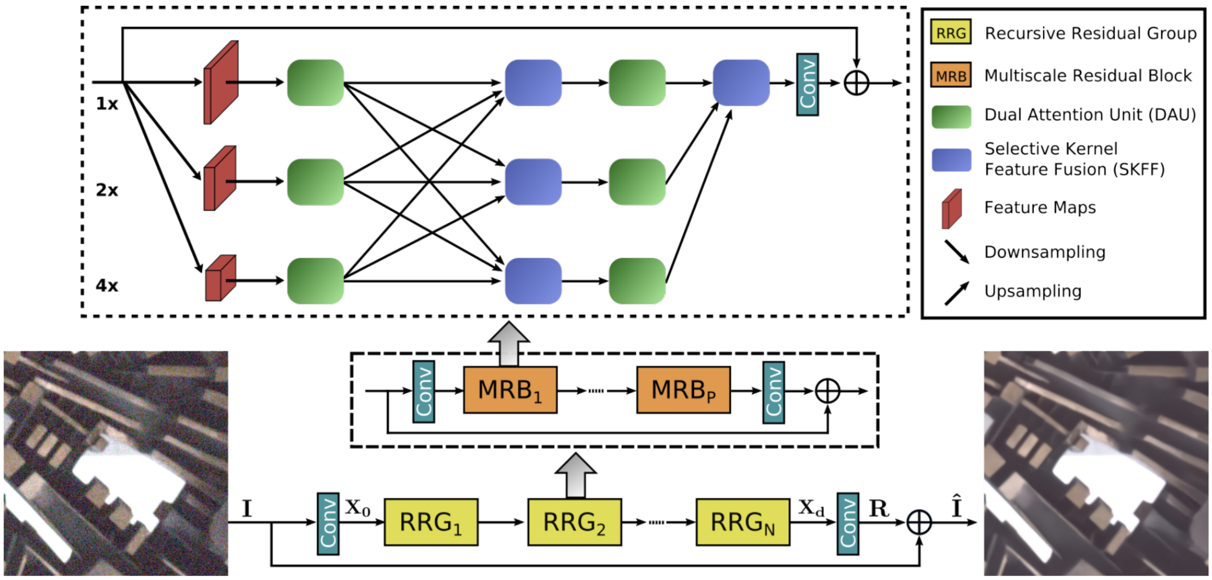

Figure 1: An overview of the MIRNet architecture.

3.2 Multi-Scale Residual Block (MSRB)

The MSRB is the fundamental building block of MIRNet, enabling the network to process and fuse features at multiple scales.

3.2.1 Structure

An MSRB consists of:

- **Parallel Convolutional Streams**: Operate at different resolutions (full, half, and quarter resolutions).
- **Downsampling and Upsampling Modules**: Adjust the resolution of feature maps.
- **Dual Attention Units**: Enhance features within each stream.
- **Selective Kernel Feature Fusion**: Combine features from different scales.

3.2.2 Workflow

1. **Input Features**: The input feature map is fed into the MSRB.
2. **Parallel Processing**:
    - **Downsampling**: The input is downsampled to create lower-resolution versions using strided convolutions or pooling.
    - **Convolutions and DAUs**: Each resolution stream applies convolutional layers and Dual Attention Units to extract and refine features.

3. **Feature Fusion**: Use SKFF to merge features from different scales.

4. **Residual Connection**: Add the aggregated features back to the input feature map.

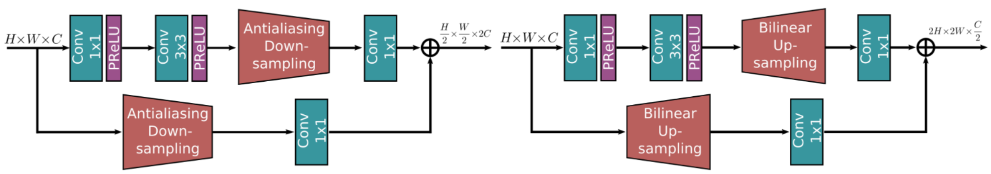

(a) Downsampling module (b) Upsampling module

Figure 2: Structure of the Multi-Scale Residual Block (MSRB).

3.3 Dual Attention Unit (DAU)

The DAU refines feature maps using two attention mechanisms: Channel Attention and Spatial Attention.

3.3.1 Channel Attention (CA)

Emphasizes 'what' to focus on across feature channels.

Process

1. **Global Pooling**: Compute channel-wise statistics using global average pooling.

2. **Fully Connected Layers**: Capture inter-channel relationships with compression and expansion layers.

3. **Scaling**: Generate channel-wise weights to re-scale the feature map.

3.3.2 Spatial Attention (SA)

Emphasizes 'where' to focus within spatial dimensions.

Process

1. **Pooling Operations**: Apply max and average pooling across channels to generate two 2D maps.

2. **Concatenation and Convolution**: Concatenate the pooled maps and apply convolution to generate spatial attention weights.

3. **Scaling**: Re-scale the feature map spatially.

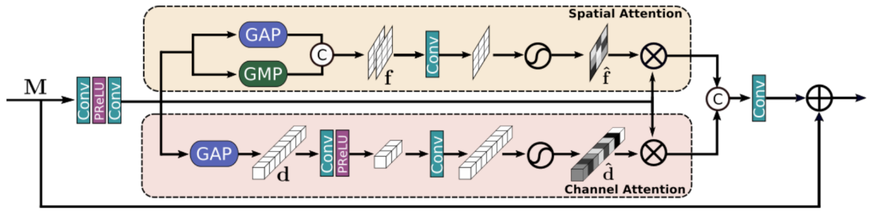

Figure 3: Structure of the Dual Attention Unit (DAU).

3.4 Selective Kernel Feature Fusion (SKFF)

SKFF dynamically fuses features from multiple scales, allowing the network to adaptively select the most informative features at each location.

3.4.1 Workflow

1. **Input Feature Maps**: Receive feature maps from different resolution streams.
2. **Global Context Embedding**:
   - Apply global average pooling to each feature map to generate scale-specific descriptors.
3. **Adaptive Weights Generation**:
   - Concatenate global descriptors.
   - Pass through fully connected layers with softmax activation to generate scale-wise attention weights.
4. **Feature Recalibration**:
   - Multiply each scale's feature map by its corresponding attention weight.
5. **Feature Aggregation**:
   - Sum the recalibrated feature maps to produce the fused output.

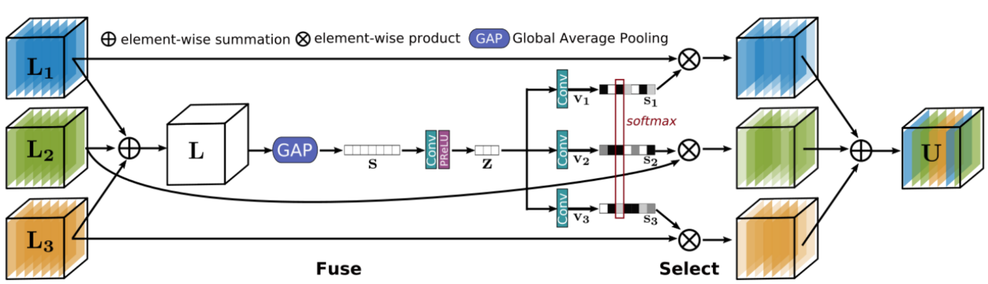

Figure 4: Structure of the Selective Kernel Feature Fusion (SKFF) mechanism.

3.5 Recursive Residual Groups (RRG)

RRGs stack multiple MSRBs recursively, aiding in deep network training through efficient gradient flow.

3.5.1 Structure

- **Sequential MSRBs**: A series of MSRBs connected in sequence.
- **Long Skip Connections**: Residual connections that bypass the entire group.

3.5.2 Workflow

1. **Input Feature Map**: Provided to the first MSRB in the group.
2. **Processing**: Pass through a series of MSRBs.
3. **Residual Connection**: Add the input feature map to the output of the final MSRB.
4. **Output Feature Map**: Passed to the next RRG or the reconstruction layers.

3.6 Overall Network Architecture

3.6.1 High-Level Structure

1. **Initial Convolutional Layer**:
   - Extracts initial features from the input image using a convolutional layer.
2. **Recursive Residual Groups**:
   - Stack multiple RRGs to deepen the network.
3. **Feature Reconstruction Layer**:
   - A convolutional layer that maps the refined features back to the image space, producing the residual image.
4. **Global Residual Connection**:
   - Adds the input image to the residual image to obtain the final enhanced output.

3.6.2 Mathematical Formulation

$$
\begin{aligned}
F_0 &= \text{Conv2D}(I) \\
F_r &= \text{RRG}(F_{r-1}), \quad r = 1, 2, \ldots, R \\
F_{\text{rec}} &= \text{Conv2D}(F_R) \\
O &= I + F_{\text{rec}}
\end{aligned}
$$

Where:

- I is the input image.
- F_0 is the initial feature map.
- F_r is the output from the r^{th} RRG.
- F_{rec} is the reconstructed features.
- O is the output (enhanced) image.

4 Implementation Details

4.1 Dataset Preparation

The **LoL Dataset** is used for training and evaluation, which contains paired low-light and well-exposed images.

- **Data Augmentation**:
  - Random cropping to generate 128×128 patches.
  - Horizontal and vertical flips, rotations.
- **Normalization**:
  - Pixel values scaled to $[0, 1]$.

4.2 Model Construction

```
import tensorflow as tf
from tensorflow.keras import layers, Model

def dual_attention_unit(x, channels):
    # Channel Attention
    avg_pool = tf.reduce_mean(x, axis=[1, 2], keepdims=True)
    max_pool = tf.reduce_max(x, axis=[1, 2], keepdims=True)
    concat = tf.concat([avg_pool, max_pool], axis=-1)
    mlp = layers.Dense(channels // 8, activation='relu')(concat)
    mlp = layers.Dense(channels, activation='sigmoid')(mlp)
    channel_attention = x * mlp

    # Spatial Attention
    avg_pool_sp = tf.reduce_mean(channel_attention, axis=-1, keepdims=True)
    max_pool_sp = tf.reduce_max(channel_attention, axis=-1, keepdims=True)
    concat_sp = tf.concat([avg_pool_sp, max_pool_sp], axis=-1)
    conv_sp = layers.Conv2D(1, kernel_size=7, padding='same', activation='sigmoid')(conc
    spatial_attention = channel_attention * conv_sp
```

```
    return spatial_attention

def multi_scale_residual_block(x, channels):
    # Assuming downsampled and upsampled features are obtained
    # This is a simplified representation
    # Apply DAU and SKFF within this block as per the architecture
    return x

def MIRNet(input_shape):
    inputs = layers.Input(shape=input_shape)
    x = layers.Conv2D(64, (3, 3), padding='same', activation='relu')(inputs)

    for _ in range(num_recursive_groups):
        res = x
        for _ in range(num_msrbs):
            x = multi_scale_residual_block(x, 64)
        x = layers.add([x, res])

    x = layers.Conv2D(3, (3, 3), padding='same')(x)
    outputs = layers.add([x, inputs])
    model = Model(inputs, outputs)
    return model

model = MIRNet(input_shape=(128, 128, 3))
```

4.3 Training Configuration

- **Loss Function**: Mean Squared Error (MSE) between the enhanced and ground truth images.
- **Optimizer**: Adam optimizer with an appropriate learning rate.
- **Metrics**: Peak Signal-to-Noise Ratio (PSNR) and Structural Similarity Index Measure (SSIM).
- **Batch Size**: 16
- **Epochs**: 200

5 Results and Evaluation

5.1 Quantitative Results

Metric	Value
PSNR	**28.95 dB**
SSIM	**0.870**

Table 1: Performance metrics on the LoL Dataset.

5.2 Qualitative Results

Figure 5: Comparison of low-light enhancement results (Example 1).

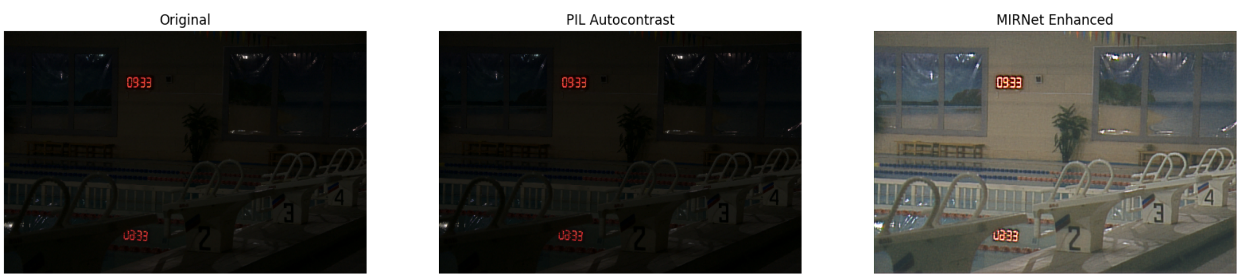

Figure 6: Comparison of low-light enhancement results (Example 2).

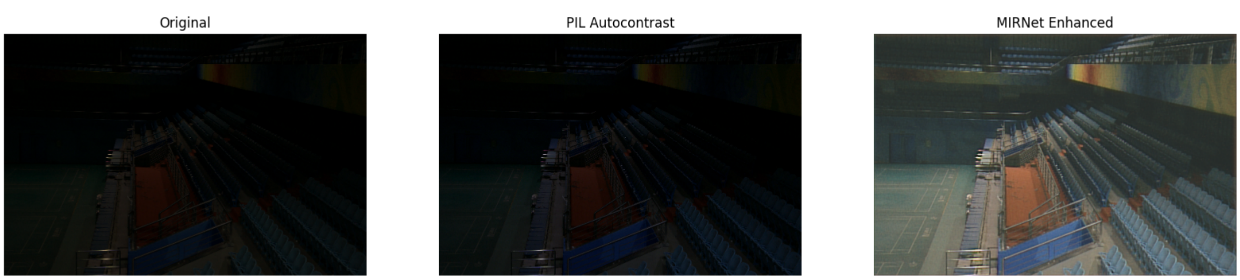

Figure 7: Comparison of low-light enhancement results (Example 2).

The enhanced images exhibit improved brightness, color fidelity, and detail preservation.

6 Discussion

6.1 Advantages

- **High-Quality Enhancement**: Produces visually pleasing results.
- **Efficient Learning**: Recursive residual design aids in training.
- **Versatility**: Can be adapted for other restoration tasks.

6.2 Limitations

- **Computational Complexity**: Increased due to multi-scale and attention mechanisms.
- **Dataset Dependency**: Performance may vary with different datasets.

6.3 Future Work

- Optimizing the model for faster inference.
- Extending to handle other types of image degradation.
- Exploring unsupervised or semi-supervised learning approaches.

7 Conclusion

The MIRNet model effectively addresses low-light image enhancement by capturing and integrating information across multiple scales and applying attention mechanisms to refine feature representations. Quantitative and qualitative results highlight the model's potential for practical applications in image processing.

8 Detailed Explanation of the MIRNet Architecture

8.1 Mathematical Formulations

8.1.1 Channel Attention (CA)

Given an input feature map $F \in R^{H \times W \times C}$:

Squeeze Operation Compute the channel-wise global average:

$$s_c = \frac{1}{H \times W} \sum_{i=1}^{H} \sum_{j=1}^{W} F_{i,j,c} \tag{1}$$

Excitation Operation Pass s through fully connected layers:

$$z = \sigma \left(W_2 \cdot \text{ReLU}(W_1 \cdot s) \right) \tag{2}$$

Re-Scaling Multiply the input feature map by the channel attention weights:

$$F_{\text{CA}} = F \odot z \tag{3}$$

8.1.2 Spatial Attention (SA)

Given F_{CA}:

Pooling Compute max and average pooling:

$$F_{\max} = \text{MaxPool}(F_{\text{CA}}, \text{axis} = C) \tag{4}$$

$$F_{\text{avg}} = \text{AvgPool}(F_{\text{CA}}, \text{axis} = C) \tag{5}$$

Concatenation and Convolution Generate spatial attention map:

$$M = \sigma \left(\text{Conv2D}(\text{Concat}[F_{\max}, F_{\text{avg}}]) \right) \tag{6}$$

Re-Scaling Multiply the feature map by the spatial attention map:

$$F_{\text{SA}} = F_{\text{CA}} \odot M \tag{7}$$

8.2 Selective Kernel Feature Fusion (SKFF)

Global Context Embedding For each scale s:

$$g_s = \frac{1}{H_s \times W_s} \sum_{i=1}^{H_s} \sum_{j=1}^{W_s} F_s(i, j) \tag{8}$$

Adaptive Weights Generation Concatenate global descriptors:

$$g = \text{Concat}[g_1, g_2, \dots, g_S] \tag{9}$$

Generate attention weights using fully connected layers with softmax:

$$\alpha = \text{Softmax} \left(W_2 \cdot \text{ReLU}(W_1 \cdot g) \right) \tag{10}$$

Feature Recalibration and Aggregation Recalibrate and sum features:

$$F_{\text{fused}} = \sum_{s=1}^{S} \alpha_s \cdot F_s \tag{11}$$

8.3 Advantages of MIRNet Architecture

- **Enhanced Feature Representation**:
  - Multi-scale processing captures diverse features pertinent to different aspects of the image.
- **Attention Mechanisms**:
  - Focus on critical features and regions, improving the enhancement quality.
- **Efficient Training**:
  - Residual connections facilitate gradient flow, allowing for deeper networks without vanishing gradients.
- **Dynamic Adaptation**:
  - SKFF enables the network to adaptively merge features, making it robust to varying image conditions.

8.4 Potential Extensions

- **Other Restoration Tasks**:
  - The architecture can be adapted for denoising, super-resolution, deblurring, etc.
- **Feature Scaling**:
  - Experimenting with different scales could capture even more diverse features.
- **Advanced Attention Mechanisms**:
  - Incorporating transformer-based attention for even more powerful feature representation.

9 Conclusion

The MIRNet model presents a sophisticated approach to low-light image enhancement by integrating multi-scale feature extraction with advanced attention mechanisms. Its architecture allows for the effective fusion of contextual and detailed information, leading to high-quality image restoration results.

10 References

10.1 References

1. *Learning Enriched Features for Real Image Restoration and Enhancement.*
2. *The Retinex Theory of Color Vision.*
3. *Two Deterministic Half-Quadratic Regularization Algorithms for Computed Imaging.*